



## CSCE 421: Machine Learning


### Lecture 6

## Overview

- Brief probability review
- Logistic Regression
  - Representation and Intuition
  - Evaluation through maximum-likelihood
  - Optimization through gradient descent
  - Convexity of evaluation criterion
- Multiclass logistic regression
  - Representation (derivation based on 2-class)
  - Evaluation through cross-entropy error
- Regularization
  - Why do we need it?
  - Non-Linear Regression
  - Logistic Regression
  - How to choose the right amount of regularization

## General Probability Review


Example: Duration (sec) to answer a Multiple Choice Question



What do you observe?

## General Probability Review

Example: Duration (sec) to answer a MCQ

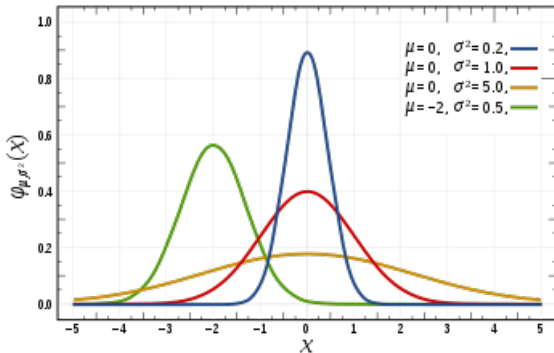


What do you observe?

It is possible that the data are generated from a Gaussian distribution, since most of the points lie in the middle, while some points are scattered to the left and the right.

## General Probability Review

### Normal distribution

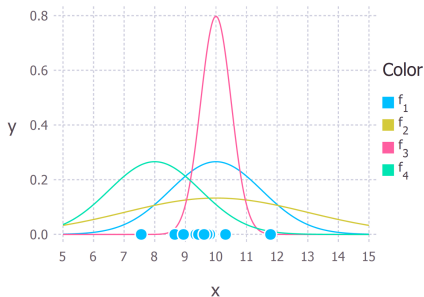


$$x \sim \mathcal{N}(0, \sigma^2) \rightarrow p(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

Mean  $\mu$ , variance  $\sigma^2$ , precision  $\tau = 1/\sigma^2$

## General Probability Review

Which model best describes the data?




$$f_1 \sim \mathcal{N}(10, 2.25), f_2 \sim \mathcal{N}(10, 9), f_3 \sim \mathcal{N}(10, 0.25), f_4 \sim \mathcal{N}(8, 2.25)$$

Is there a systematic way to find the distribution that describes “best” the data?

## General Probability Review

Which model best describes the data?



- We can calculate the distribution of observing each of the data  $x_n$   

$$p(x_n|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x_n - \mu)^2}{2\sigma^2}\right), n = 1, \dots, N$$
- Find the joint distribution of all data  $\mathcal{X} = \{x_1, \dots, x_N\}$  (likelihood)

$$p(\mathcal{X}|\mu, \sigma^2) = p(\{x_1, \dots, x_N\}|\mu, \sigma^2) = \prod_{n=1}^N p(x_n|\mu, \sigma^2)$$

$$= \prod_{n=1}^N \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x_n - \mu)^2}{2\sigma^2}\right)$$

- Find the parameters  $\mu$  and  $\sigma$  that maximize this joint distribution

## Maximum likelihood estimation

- Independent identically distributed sample  $\mathcal{X} = \{x_1, \dots, x_N\}$
- Assume all samples are drawn from the same distribution  $p(x|\theta)$
- We want to find  $\theta$  that makes sampling from  $p(x|\theta)$  as likely as possible  $\rightarrow$  **maximize likelihood**

$$l(\theta|\mathcal{X}) \equiv p(\mathcal{X}|\theta) = \prod_{n=1}^N p(x_n|\theta)$$

- **Maximum Likelihood estimator (MLE):** the parameter  $\theta^{MLE}$  that maximizes the likelihood

$$\theta^{MLE} = \max_{\theta} l(\theta|\mathcal{X})$$

- For the sake of convenience, we take the **log-likelihood**

$$\mathcal{L}(\theta|\mathcal{X}) \equiv \log p(\mathcal{X}|\theta) = \sum_{n=1}^N \log p(x_n|\theta)$$



## Maximum likelihood estimation: Examples

**Normal:** models a sample from a population with continuous values

- $X$ : Gaussian normal distributed with mean  $\mu$  and variance  $\sigma^2$
- **PDF:**  $p(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$
- **MLE estimation:** Sample  $\mathcal{X} = \{x_1, \dots, x_N\}$

$$m = \mu^{MLE} = \frac{\sum_{n=1}^N x_n}{N} \quad s^2 = (\sigma^2)^{MLE} = \frac{\sum_{n=1}^N (x_n - \mu^{MLE})^2}{N}$$

i.e. the MLE estimate for the *population mean* is the *sample mean*

**Note:** Not all continuous variables follow the normal distribution, we might have to perform statistical tests for that


## Overview

- Logistic Regression
  - Representation and Intuition
  - Evaluation through maximum-likelihood
  - Optimization through gradient descent
  - Convexity of evaluation criterion
- Multiclass logistic regression
  - Representation (derivation based on 2-class)
  - Evaluation through cross-entropy error
- Regularization
  - Why do we need it?
  - Non-Linear Regression
  - Logistic Regression
  - How to choose the right amount of regularization

## Why the sigmoid function?

Three linear models that we have seen so far

$$s = \mathbf{w}^T \mathbf{x} = \sum_{d=1}^D w_d x_d$$




With logistic regression, we can find a **soft threshold** and model **uncertainty**.

## Why the sigmoid function?

Three linear models that we have seen so far

Example of credit analysis



## Bernoulli distribution

The pdf of a single experiment asking a yes/no question

- $Y \sim \text{Bernoulli}(\theta)$ , where  $Y \in \{0, 1\}$
- $p(y|\theta) = \theta^{\mathbb{I}(y=1)}(1 - \theta)^{\mathbb{I}(y=0)} = \begin{cases} \theta & y = 1 \\ 1 - \theta & y = 0 \end{cases}$
- e.g. coin toss experiment

## Logistic Regression

Parametric classification method (not regression)

Sometimes referred as "generalization" of linear regression because

- We still compute a linear combination of feature inputs, i.e.  $\mathbf{w}^T \mathbf{x}$
- Instead of predicting a continuous output variable from  $\mathbf{w}^T \mathbf{x}$ 
  - We pass  $\mathbf{w}^T \mathbf{x}$  through a function  $\mu(\mathbf{w}^T \mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x})$

$$\mu(\eta) = \sigma(\eta) = \frac{1}{1 + e^{-\eta}}, \quad 0 \leq \mu(\eta) \leq 1$$


- The above can be considered as the parameter  $\theta$  of a Bernoulli distribution

$$p(y|\mathbf{x}, \mathbf{w}) = \text{Ber}(y|\mu(\mathbf{w}^T \mathbf{x}))$$

*The output belongs to class 1 ( $y = 1$ ) with probability  $\theta = \mu(\mathbf{w}^T \mathbf{x})$  and to class 0 ( $y = 0$ ) with probability  $1 - \theta = 1 - \mu(\mathbf{w}^T \mathbf{x})$ .*

## Why the sigmoid function?

$$\sigma(\eta) = \frac{1}{1+e^{-\eta}} = \frac{e^{\eta}}{1+e^{\eta}}$$



### Very nice properties

- Bounded between 0 and 1 ← thus interpretable as a probability
- Monotonically increasing ← thus can be used for classification rules
  - $\sigma(\eta) > 0.5$ , positive class ( $y=1$ )
  - $\sigma(\eta) \leq 0.5$ , positive class ( $y=0$ )
- Nice computational properties for optimizing criterion function

## Logistic Regression: Representation

### Setup for two classes

- **Input:**  $\mathbf{x} \in \mathbb{R}^D$
- **Output:**  $y \in \{0, 1\}$
- **Training data:**  $\mathcal{D}^{train} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$
- **Model:**


$$p(y = 1 | \mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x}), \quad \sigma(\eta) = \frac{1}{1 + e^{-\eta}}$$

$$y = \begin{cases} 1, & p(y = 1 | \mathbf{x}, \mathbf{w}) > 0.5 \\ 0, & \text{otherwise} \end{cases}$$

- **Model parameters:** weights  $\mathbf{w}$



## Logistic Regression



**Classification task:** whether a student passes or not the class

**Features:** SAT scores

**Data:** SAT scores v.s. fail/pass ( $y=0/1$ ) (solid black dots)

**Logistic regression:**

- Assigns each score to “pass” probability (open red circles)
- If  $p(y = 1|x) > 0.5$ , then decides  $y(x) = 1$ . Otherwise,  $y(x) = 0$ .

## Logistic Regression: Evaluation

Data likelihood for 1 training sample

$$p(y_n|\mathbf{x}_n, \mathbf{w}) = \left\{ \begin{array}{ll} \sigma(\mathbf{w}^T \mathbf{x}_n), & y_n = 1 \\ 1 - \sigma(\mathbf{w}^T \mathbf{x}_n), & y_n = 0 \end{array} \right\} = [\sigma(\mathbf{w}^T \mathbf{x}_n)]^{y_n} [1 - \sigma(\mathbf{w}^T \mathbf{x}_n)]^{1-y_n}$$

Data likelihood for all training data

$$L(\mathcal{D}|\mathbf{w}) = \prod_{n=1}^N p(y_n|\mathbf{x}_n, \mathbf{w}) = \prod_{n=1}^N [\sigma(\mathbf{w}^T \mathbf{x}_n)]^{y_n} [1 - \sigma(\mathbf{w}^T \mathbf{x}_n)]^{1-y_n}$$

Cross-entropy error (negative log-likelihood)

$$\begin{aligned} \mathcal{E}(\mathbf{w}) &= -\log L(\mathcal{D}|\mathbf{w}) \\ &= -\sum_{n=1}^N \{y_n \log [\sigma(\mathbf{w}^T \mathbf{x}_n)] + (1 - y_n) \log [1 - \sigma(\mathbf{w}^T \mathbf{x}_n)]\} \end{aligned}$$

## Logistic Regression: Optimization

Cross-entropy error (negative log-likelihood)

$$\mathcal{E}(\mathbf{w}) = - \sum_{n=1}^N \{ y_n \log [\sigma(\mathbf{w}^T \mathbf{x}_n)] + (1 - y_n) \log [1 - \sigma(\mathbf{w}^T \mathbf{x}_n)] \}$$

How to find the weights  $\mathbf{w}$  of the logistic regression?

We can maximize data likelihood or minimize cross-entropy error

$$\mathbf{w}^* = \min_{\mathbf{w}} \mathcal{E}(\mathbf{w})$$

No closed-form solution  $\rightarrow$  approximate methods, e.g. **Gradient Descent**.

$$\mathbf{w} := \mathbf{w} - \alpha(k) \cdot \nabla \mathcal{E}(\mathbf{w}), \quad \frac{\partial \mathcal{E}(\mathbf{w})}{\partial w_d} = \sum_{n=1}^N \underbrace{(\sigma(\mathbf{w}^T \mathbf{x}_n) - y_n)}_{\text{error}} x_{nd}$$

$\mathcal{E}(\mathbf{w})$  is convex, i.e. has a global minimum (**positive definite Hessian**).

## Overview

- Logistic Regression
  - Representation and Intuition
  - Evaluation through maximum-likelihood
  - Optimization through gradient descent
  - Convexity of evaluation criterion
- Multiclass logistic regression
  - Representation (derivation based on 2-class)
  - Evaluation through cross-entropy error
- Regularization
  - Why do we need it?Non-
  - Non-Linear Regression
  - Logistic Regression
  - How to choose the right amount of regularization

## Multi-class logistic regression

- Suppose we need to predict multiple classes/outcomes  $1, \dots, C$ 
  - weather prediction: rainy, cloudy, shiny
  - optical digit/character recognition: 0-9 or 'a'-'z'

- 2-class: probability of  $\mathbf{x}$  belonging to class 1

$$p(y = 1|\mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x}), \quad \sigma(\eta) = \frac{1}{1+e^{-\eta}} = \frac{e^{\eta}}{1+e^{\eta}}$$

- How could we generalize to  $C$  classes?

- One way could be  $p(y = c|\mathbf{x}, \mathbf{w}_c) = \sigma(\mathbf{w}_c^T \mathbf{x}) = \frac{e^{\mathbf{w}_c^T \mathbf{x}}}{1+e^{\mathbf{w}_c^T \mathbf{x}}}$
- This would not work, because each  $p(y = c|\mathbf{x}, \mathbf{w}_c) \in [0, 1]$  independently
- And we need  $\sum_{c=1}^C p(y = c|\mathbf{x}, \mathbf{w}_c) \in [0, 1]$

- But we can do the following (**softmax function** or **conditional logit model**)

$$p(y = c|\mathbf{x}, \mathbf{w}_c) = \frac{e^{\mathbf{w}_c^T \mathbf{x}}}{\sum_{c=1}^C e^{\mathbf{w}_c^T \mathbf{x}}} = \frac{e^{\mathbf{w}_c^T \mathbf{x}}}{e^{\mathbf{w}_1^T \mathbf{x}} + \dots + e^{\mathbf{w}_C^T \mathbf{x}}}$$

$$\sum_{c=1}^C p(y = c|\mathbf{x}, \mathbf{w}_c) = 1$$

## Multi-class logistic regression

- **Input:**  $\mathbf{x} \in \mathbb{R}^D$
- **Output:**  $y \in \{1, 2, \dots, C\}$
- **Training data:**  $\mathcal{D}^{train} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$
- **Model:**

$$p(y = c | \mathbf{x}, \mathbf{w}_c) = \frac{e^{\mathbf{w}_c^T \mathbf{x}}}{\sum_{c=1}^C e^{\mathbf{w}_c^T \mathbf{x}}}$$

$$y = \arg \max_{c=1, \dots, C} p(y = c | \mathbf{x}, \mathbf{w}_c)$$

- **Model parameters:** weights  $\mathbf{w}_1, \dots, \mathbf{w}_C$

## Multi-class logistic regression

Binary logistic regression is a special case of multi-class

From  $p(y = c|\mathbf{x}, \mathbf{w}_c) = \frac{e^{\mathbf{w}_c^T \mathbf{x}}}{\sum_{c=1}^C e^{\mathbf{w}_c^T \mathbf{x}}}$  for  $c = \{0, 1\}$ , we get

$$p(y = 1|\mathbf{x}, \mathbf{w}_c) = \frac{e^{\mathbf{w}_1^T \mathbf{x}}}{e^{\mathbf{w}_0^T \mathbf{x}} + e^{\mathbf{w}_1^T \mathbf{x}}} = \frac{1}{e^{\mathbf{w}_0^T \mathbf{x} - \mathbf{w}_1^T \mathbf{x}} + 1} = \frac{1}{1 + e^{(\mathbf{w}_0 - \mathbf{w}_1)^T \mathbf{x}}}$$

Same as  $p(y = 1|\mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x})$  with  $\mathbf{w} = \mathbf{w}_0 - \mathbf{w}_1$

## Multi-class logistic regression: Optimization

### Discriminative Approach

- We will change  $y_n \in \mathbb{R}$  to a C-dimensional vector

$$\mathbf{y}_n = [y_{n1}, \dots, y_{nC}]^T \in \mathbb{R}^C$$

$$y_{nc} = \begin{cases} 1, & \text{if } y_n = c \\ 0, & \text{otherwise} \end{cases}$$

e.g. if  $y_n = 3$  then  $\mathbf{y}_n = [0, 0, 1, 0, \dots, 0]^T \in \mathbb{R}^C$

- We will maximize the likelihood

$$\begin{aligned} L(\mathcal{D}|\mathbf{w}_1, \dots, \mathbf{w}_C) &= \prod_{n=1}^N p(\mathbf{y}_n|\mathbf{x}_n) \\ &= \prod_{n=1}^N (p(y_{n1} = 1|\mathbf{w}_1, \dots, \mathbf{w}_C)^{y_{n1}} \dots p(y_{nC} = 1|\mathbf{w}_1, \dots, \mathbf{w}_C)^{y_{nC}}) \end{aligned}$$



## Multi-class logistic regression: Optimization

### Data-likelihood

$$\begin{aligned}
 L(\mathcal{D}|\mathbf{w}_1, \dots, \mathbf{w}_C) &= \prod_{n=1}^N p(y_n|\mathbf{x}_n) \\
 &= \prod_{n=1}^N (p(y_{n1} = 1|\mathbf{w}_1, \dots, \mathbf{w}_C)^{y_{n1}} \dots p(y_{nC} = 1|\mathbf{w}_1, \dots, \mathbf{w}_C)^{y_{nC}}) \\
 &= \prod_{n=1}^N \prod_{c=1}^C p(y_{nc} = 1|\mathbf{w}_1, \dots, \mathbf{w}_C)^{y_{nc}}
 \end{aligned}$$

### Cross-entropy error

$$\mathcal{E}(\mathbf{w}_1, \dots, \mathbf{w}_C) = - \sum_{n=1}^N \sum_{c=1}^C y_{nc} \log p(y_{nc} = 1|\mathbf{w}_1, \dots, \mathbf{w}_C)$$

## Multi-class logistic regression: Optimization

### Cross-entropy error

$$\mathcal{E}(\mathbf{w}_1, \dots, \mathbf{w}_C) = - \sum_{n=1}^N \sum_{c=1}^C y_{nc} \log p(y_{nc} = 1 | \mathbf{w}_1, \dots, \mathbf{w}_C)$$

- Optimization with gradient descent, convex function
- Computational details are out of scope
- But the gradient vector w.r.t. each weight  $\mathbf{w}_c$  looks like this

$$\nabla \mathcal{E}_{\mathbf{w}_c} = \sum_{n=1}^N \underbrace{[p(y_{nc} = 1 | \mathbf{w}_1, \dots, \mathbf{w}_C) - y_{nc}]}_{\text{error for class } c} \mathbf{x}_n$$





- Similar to binary logistic regression  $\rightarrow$  General property of exponential family distributions

## Overview


- Logistic Regression
  - Representation and Intuition
  - Evaluation through maximum-likelihood
  - Optimization through gradient descent
  - Convexity of evaluation criterion
- Multiclass logistic regression
  - Representation (derivation based on 2-class)
  - Evaluation through cross-entropy error
- Regularization
  - Why do we need it?
  - Non-Linear Regression
  - Logistic Regression
  - How to choose the right amount of regularization

## Overfitting

**Example:** Non-linear regression  $y = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M$   
 Samples from a sine function  $x_i = \sin(t_i)$ ,  $t_i \sim \text{Uniform}(0, 2\pi)$






	$M = 0$	$M = 1$	$M = 3$	$M = 9$
$w_0$	0.19	0.82	0.31	0.35
$w_1$		-1.27	7.99	232.37
$w_2$			-25.43	-5321.83
$w_3$			17.37	48568.31
$w_4$				-231639.30
$w_5$				640042.26
$w_6$				-1061800.52
$w_7$				1042400.18
$w_8$				-557682.99
$w_9$				125201.43



As model becomes more complex, performance on training keeps improving while on test data improve first and deteriorate later.  
 The larger a coefficient  $w_i$ , the easier for the model to "swing" in that dimension, increasing chance to fit more noise.

## How can we avoid overfitting?

**One solution:** Use more training data




What if we don't have a lot of data?

**Another solution:** Use less features (e.g. feature selection algorithms)


Intuitively, this will reduce the complexity of the model, therefore it is likely to result in less overfitting.

## How can we avoid overfitting?

A more general solution: Regularization



$$y = w_0 + w_1x + w_2x^2$$



$$y = w_0 + w_1x + w_2x^2 + \dots + w_9x^9$$

How about penalizing and making small  $w_3, \dots, w_9$ ?

The cost function to be minimized would become:

$$J(\mathbf{w}) = \text{RSS}(\mathbf{w}) + w_3^2 + \dots + w_9^2$$

But we may not know in advance which parameters we want to penalize

→ So we can penalize them all

## How can we avoid overfitting?

### A more general solution: Regularization

Suppose we have a learning model whose evaluation criterion  $EC(\mathbf{w})$  we want to optimize with respect to weights  $\mathbf{w} = [w_1, \dots, w_D]^T$

- $J(\mathbf{w}) = EC(\mathbf{w}) + \lambda \sum_{d=1}^D w_d^2 = EC(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$

→ l2-norm regularization

- $J(\mathbf{w}) = EC(\mathbf{w}) + \frac{\lambda}{N} \sum_{d=1}^D w_d^2$

(as #data  $N$  increases, we need to worry less about overfitting)

- $J(\mathbf{w}) = EC(\mathbf{w}) + \lambda \sum_{d=1}^D \|w_d\| = EC(\mathbf{w}) + \lambda \|\mathbf{w}\|$

→ l1-norm regularization

Evaluation criterion  $EC(\mathbf{w})$  can be RSS or log-likelihood for linear regression, negative cross-entropy for logistic regression, etc.

$\lambda \geq 0$  is the model complexity penalty

## Overview

- Logistic Regression
  - Representation and Intuition
  - Evaluation through maximum-likelihood
  - Optimization through gradient descent
  - Convexity of evaluation criterion
- Multiclass logistic regression
  - Representation (derivation based on 2-class)
  - Evaluation through cross-entropy error
- Regularization
  - Why do we need it?
  - Non-Linear Regression
  - Logistic Regression
  - How to choose the right amount of regularization



## Regularization for Non-Linear Regression

### $\ell_2$ -norm regularization

Linear:  $J(\mathbf{w}) = \text{RSS}(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2 = (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$

Non-linear:  $J(\mathbf{w}) = \text{RSS}(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2 = (\mathbf{y} - \Phi\mathbf{w})^T (\mathbf{y} - \Phi\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$

### Closed-form solution:

Linear:  $\mathbf{w}^* = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}_{D \times D})^{-1} \mathbf{X}^T \mathbf{y}$


Non-linear:  $\mathbf{w}^* = (\Phi^T \Phi + \lambda \mathbf{I}_{D \times D})^{-1} \Phi^T \mathbf{y}$

The above reduces to ordinary least squares (OLS) solution when  $\lambda = 0$   
(see handout for derivation)


## Regularization for Non-Linear Regression

**Question:** Assume a set of samples generated from a sine function  $x_i = \sin(t_i)$  (green line), modeled with **regularized** non-linear regression  $y = w_0 + w_1x + \dots + w_9x^9$ . How does the resulting model (red line) look as we increase the amount of regularization  $\lambda$ ?


A)



B)






C)



## Regularization for Non-Linear Regression

**Question:** Assume a set of samples generated from a sine function  $x_i = \sin(t_i)$  (green line), modeled with **regularized** non-linear regression  $y = w_0 + w_1x + \dots + w_9x^9$ . How does the resulting model (red line) look as we increase the amount of regularization  $\lambda$ ?


The correct answer is A




Overfitting is reduced with the help of increasing regularizers

## Regularization for Non-Linear Regression



$\lambda = 0$  : no regularization



$\lambda = e^{-18}$



$\lambda = 1$



	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
$w_0$	0.35	0.35	0.13
$w_1$	232.37	4.74	-0.05
$w_2$	-5321.83	-0.77	-0.06
$w_3$	48568.31	-31.97	-0.06
$w_4$	-231639.30	-3.89	-0.03
$w_5$	640042.26	55.28	-0.02
$w_6$	-1061800.52	41.32	-0.01
$w_7$	1042400.18	-45.95	-0.00
$w_8$	-557682.99	-91.53	0.00
$w_9$	125201.43	72.68	0.01

For a complex model ( $M = 9$ ), training error increases with increasing regularization.

## Overview

- Logistic Regression
  - Representation and Intuition
  - Evaluation through maximum-likelihood
  - Optimization through gradient descent
  - Convexity of evaluation criterion
- Multiclass logistic regression
  - Representation (derivation based on 2-class)
  - Evaluation through cross-entropy error
- Regularization
  - Why do we need it?
  - Non-Linear Regression
  - Logistic Regression
  - How to choose the right amount of regularization

## Regularization for Logistic Regression

### $\ell_2$ -norm regularization

$$\mathcal{E}(\mathbf{w}) = - \sum_{n=1}^N \left\{ y_n \log [\sigma(\mathbf{w}^T \mathbf{x}_n)] + (1 - y_n) \log [1 - \sigma(\mathbf{w}^T \mathbf{x}_n)] \right\} + \lambda \|\mathbf{w}\|_2^2$$

$$\nabla \mathcal{E}(\mathbf{w}) = \sum_{n=1}^N (\sigma(\mathbf{w}^T \mathbf{x}_n) - y_n) \mathbf{x}_n + 2\lambda \mathbf{w}$$

$$\mathbf{H} = \sum_{n=1}^N \underbrace{\sigma(\mathbf{w}^T \mathbf{x}_n)}_{\in [0,1]} \cdot \underbrace{(1 - \sigma(\mathbf{w}^T \mathbf{x}_n))}_{\in [0,1]} \cdot \underbrace{(\mathbf{x}_n \cdot \mathbf{x}_n^T)}_{\in \mathcal{R}^{D \times D}} + \lambda \mathbf{I}_{D \times D}$$

(see handout for derivations)

## Overview

- Logistic Regression
  - Representation and Intuition
  - Evaluation through maximum-likelihood
  - Optimization through gradient descent
  - Convexity of evaluation criterion
- Multiclass logistic regression
  - Representation (derivation based on 2-class)
  - Evaluation through cross-entropy error
- Regularization
  - Why do we need it?
  - Non-Linear Regression
  - Logistic Regression
  - How to choose the right amount of regularization


## Overview

- Logistic Regression
  - Representation and Intuition
  - Evaluation through maximum-likelihood
  - Optimization through gradient descent
  - Convexity of evaluation criterion
- Multiclass logistic regression
  - Representation (derivation based on 2-class)
  - Evaluation through cross-entropy error
- Regularization
  - Why do we need it?
  - Non-Linear Regression
  - Logistic Regression
  - How to choose the right amount of regularization



## How to choose the right amount of regularization?

- We **cannot** tune  $\lambda$  on the train set. **Why?**
- $\lambda$  is a **hyper-parameter** and we can tune it by:
  - keeping out a hold-out-set independent of train and test sets
  - doing cross-validation
  - similar procedure to choosing  $K$  for K-NN



## Recipe for cross-validation for choosing $\lambda$

- Split train data into  $S$  equal parts, each noted as  $\mathcal{D}_s^{train}$ ,  $s = 1, \dots, S$
- For each hyperparameter value (e.g.  $\lambda = 10^{-5}, 10^{-4}, \dots$ )
  - For each  $s = 1, \dots, S$ 
    - Train model using  $\mathcal{D}^{train} \setminus \mathcal{D}_s^{train}$
    - Evaluate model performance (noted as  $E_s$ ) on  $\mathcal{D}_s^{train}$
  - Compute average performance for current hyperparameter
 
$$E = \frac{1}{S} \sum_{s=1}^S E_s$$
- Chose the hyperparameter corresponding to best average performance  $E$
- Use the best hyperparameter to train on a model using all  $\mathcal{D}^{train}$
- Evaluate the last model on  $\mathcal{D}^{test}$

## What have we learnt so far

### Logistic Regression

- Linear combination of input features  $\mathbf{w}^T \mathbf{x}$
- Transform through sigmoid function  $\sigma(\mathbf{w}^T \mathbf{x}) \rightarrow$  interpretable as probability
- Decision rule based on whether  $\sigma(\mathbf{w}^T \mathbf{x}) \leq 0.5$
- Evaluation through data likelihood, or cross-entropy error

$$\mathcal{E}(\mathbf{w}) = - \sum_{n=1}^N \{ y_n \log [\sigma(\mathbf{w}^T \mathbf{x}_n)] + (1 - y_n) \log [1 - \sigma(\mathbf{w}^T \mathbf{x}_n)] \}$$

- Optimization through gradient descent

## What have we learnt so far

### Multinomial Regression

- Conditional logit model:  $p(y = c | \mathbf{x}, \mathbf{w}_c) = \frac{e^{\mathbf{w}_c^T \mathbf{x}}}{\sum_{c=1}^C e^{\mathbf{w}_c^T \mathbf{x}}}$
- Similar to 2-class logistic regression
  - compute negative cross-entropy and perform gradient descent

### Regularization

- Method to avoid overfitting
- Penalize large weights with l1 or l2-norm regularization
 
$$J(\mathbf{w}) = EC(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$$

Readings: Alpaydin 10.7; Abu-Mostafa 3