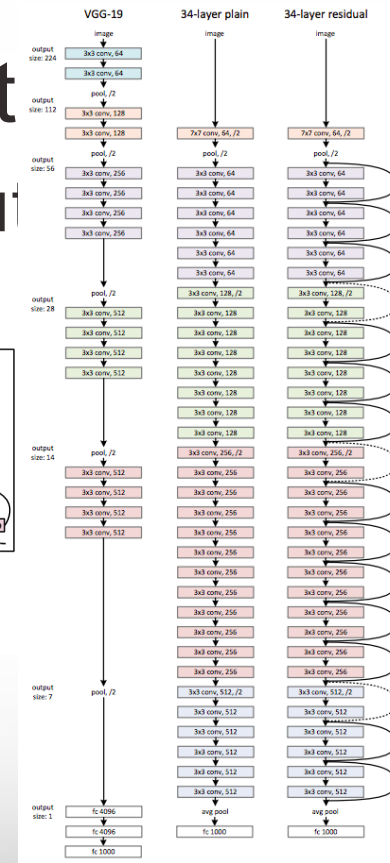
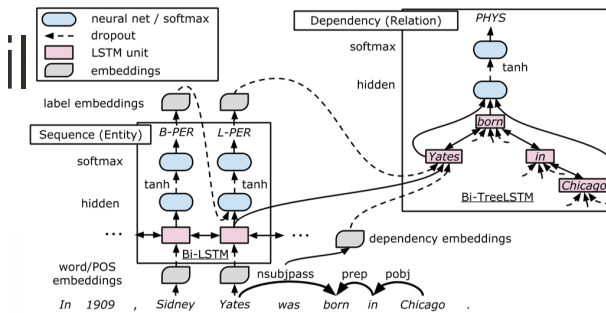



A Tutorial on PyTorch

Speaker : Zhenyu Wu




- Complicated DL architectures
 - Easily build big computational graphs




Why a DL Library is Necessary?

- Run it all efficiently on GPU

CPU (Multiple Cores)



GPU (Hundreds of Cores)



Popular Deep Learning Libraries

Caffe

Microsoft
CNTK

dmlc
mxnet



Caffe2


TensorFlow


Chainer

theano

 torch



PYTORCH

 PaddlePaddle

Keras vs PyTorch vs TensorFlow



Keras:

- High-level API
- On top of TensorFlow, CNTK, or Theano
- Easy to use
- Less flexible


PYTORCH



PyTorch (Facebook) vs TensorFlow (Google):

- Low-level API
- Flexible to write any complex models (research)
- Dynamic vs Static computation graphs

Percent of ML papers that mention...



Performance Comparison

Test VGG on CIFAR-10

DL Library	Test Accuracy (%)	Training Time (s)
Caffe2	79	149
MXNet	77	149
Gluon	77	157
CNTK	78	166
PyTorch	78	168
Tensorflow	78	173
Keras(CNTK)	78	200

Test LSTM on IMDB
GPU-accelerated LSTM

DL Library	Test Accuracy (%)	Training Time (s)
Keras(CNTK)	86	223
Tensorflow	86	79
Pytorch	87	36
MXNet	88	12

Why



- Numpy-like Tensor Calculation
 - `numpy.reshape()` → `torch.view()`
 - `numpy.concatenate()` → `torch.cat()`
 - `numpy.dot()` → `torch.dot()`
 - Support slicing, indexing, broadcasting
- Powerful tensor calculation with GPU support
- Flexible auto-differentiation & auto-grad system
- Dynamic Computation Graph (suitable for NLP research)
- Good community support and documentation
 - Latest deep learning models: GAN, VGG, ResNet, seq2seq

PyTorch as A Tensor Library

- Tensor operations: slicing, indexing, math operations, linear algebra, reductions
 - CPU & GPU
 - Fast! (comparison on speed of **matrix multiplication**)

$$\mathbf{M} * \mathbf{M} * \mathbf{M} \quad M \in \mathbb{R}^{1000 \times 1000}$$

Numpy


```
In [2]: M = numpy.random.randn(1000,1000)
In [3]: timeit -n 500 M.dot(M).dot(M)
500 loops, best of 3: 30.7 ms per loop
```

PyTorch

```
In [4]: N = torch.randn(1000,1000).cuda()
In [5]: timeit -n 500 N.mm(N).mm(N)
500 loops, best of 3: 474 µs per loop
```




Dynamic computation graph



(a) *Define-and-Run*: existing approach


Static graph



(b) *Define-by-Run*: new approach

Dynamic graph

Using GPU with Pytorch



- Use `.cuda()` on your data (stored in Variable)
- Use `.cuda()` on your model (stored in Module)

```
model = MyModel(...)
```

```
if torch.cuda.is_available():
```

```
    X.cuda()
```

```
    model.cuda()
```

Pretrained model in Pytorch

Super easy to use pretrained models with torchvision

<https://github.com/pytorch/vision>

```
import torch
import torchvision

alexnet = torchvision.models.alexnet(pretrained=True)
vgg16 = torchvision.models.vgg16(pretrained=True)
resnet101 = torchvision.models.resnet101(pretrained=True)
```



Installation



PyTorch Build	Stable (1.3)		Preview (Nightly)		
Your OS	Linux		Mac	Windows	
Package	Conda		Pip	LibTorch	Source
Language	Python 2.7	Python 3.5	Python 3.6	Python 3.7	C++
CUDA	9.2		10.1	None	
Run this Command:	<code>conda install pytorch torchvision cudatoolkit=9.2 -c pytorch</code>				



Thanks!

